

Content Therapy

Cinema as medicine

Adaptive cinematic narratives, generated on a clinician's brief and delivered as screenlife on the patient's own device — one engine, three revenue surfaces.

“Film is a machine that generates empathy.”

– Roger Ebert · “We turn this machine into medicine.” – Stanislavsky AI

01 · What is it?	02 · Why is it important?
03 · How to implement it?	04 · Who is it for?
05 · How much is it?	06 · How long will it take?
07 · Who makes it?	08 · Pre-mortem & verified data

One narrative engine. Three commercial surfaces.

Content Therapy builds the Conductor Engine: psychological profiling → emotional-trajectory computation → adaptive multi-modal narrative delivered through messages, calls, notifications and video on the user's own device.

REVENUE · MONTHS 4–6 Entertainment ARG Interactive screenlife narratives for studios and brands. Project fees (\$150–300K per campaign) extend runway and produce the brand-name case study Series A requires.	REVENUE · MONTHS 9–12 Consumer emotional fitness Daily personalized narrative, positioned as wellness — no treatment claims. Priced against the real anchors people already pay: therapy ~\$143/session, ChatGPT \$20/mo.	POST-SERIES A Prescription digital therapeutic Adjunctive indication (insomnia first, depression second) on the Rejoyn/Somryst precedent — pursued only once entertainment + employer revenue carries the company.
---	--	---

The same engine that drives entertainment engagement enforces clinical structure: separate hard-coded optimization targets per surface — engagement metrics for ARG, validated instruments (PHQ-9, GAD-7, ISI) for the clinical loop.

The failure mode of digital mental health isn't evidence. It's experience.

19.1%

of U.S. adults report current depression — a record high (Gallup, 2026)

~\$143

average U.S. therapy session out-of-pocket (Headway, 175K+ therapists); ~1 in 3 psychologists take no insurance (APA, 2024)

55.9%

pooled completion of digital CBT-I programs (npj Digital Medicine meta-analysis, 2025) — half of users never finish

The graveyard proves the point

- ▶ Pear Therapeutics — FDA-cleared, bankrupt Apr 2023; \$1.6B valuation → assets sold for ~\$6M. [\[STAT\]](#)
- ▶ Woebot — most credible AI-therapy app, consumer shutdown Jun 30, 2025. [\[STAT\]](#)
- ▶ BetterHelp — refunded \$93M in 2023 (~8% of revenue). [\[BHB\]](#)
- ▶ Users call AI mental-health apps “detached, ignoring real-life circumstances” — an experience problem, not a content problem.

Why now

- ▶ Payment rails finally exist: CMS finalized the first billing codes for FDA-cleared digital mental-health devices (G0552–G0554), effective Jan 1, 2025; expanded again for 2026. [\[CMS\]](#)
- ▶ Market: AI-in-mental-health ~\$1.7B (2025) → \$5.1B by 2030, ~23% CAGR; mental-health apps → \$17.5B by 2030. [\[Grand View\]](#)
- ▶ Science is ready: film synchronizes >65% of viewers' cortex (Hasson, neurocinematics) [\[Hasson\]](#); expectation alone matched drug effect in Harvard's migraine RCT (Kam-Hansen 2014). [\[PubMed\]](#)
- ▶ Generative video cost collapse makes per-patient cinematic content economically possible for the first time.

A prescription for cinema — with clinical guardrails designed in.

STEP 1 Clinical brief Physician writes trigger profile, diagnosis context, emotional profile, therapeutic goal	STEP 2 Creator agent Patient's "Creator DNA" profiled through Stanislavsky-method dialogue	STEP 3 Conductor engine Therapeutic objective → Cinematic Therapeutic Score	STEP 4 Rendering AI-orchestrated video, music, voice, sound design	STEP 5 Delivery + review Device-native format; clinician measures before/after every session
---	---	--	---	---

EMDR-informed protocol — repositioned for safety

Three films (*The Safe Place* → *The Boat* → *Friday*) map to stabilization → symbolic processing → integration. Per the clinical revision (V. Kravitz, May 2026), we present this as an “EMDR-informed audiovisual trauma-recovery protocol for supervised clinical pilot testing” — never as a replacement for standard EMDR. Bilateral audio-visual elements are *cinematic analogues* of bilateral stimulation; their clinical contribution is an empirical question the pilot exists to answer.

Safety architecture

- ▶ Screening & exclusion: acute suicidality, psychosis, mania, severe dissociation, unstable substance use — excluded until stabilized
- ▶ Stop protocol + grounding sequence hard-coded; clinician supervises every phase transition
- ▶ Measured, not declared: PCL-5, SUD/VOC pre-post, sleep metrics, adverse-event monitoring; ISI + PHQ-9/GAD-7 for insomnia/depression tracks
- ▶ RCT design ready: 4-arm, N=120 (standard care / generic video / personalized / personalized + ritual context)

Science base: placebo neurobiology — μ -opioid mediation (Zubieta 2005), fMRI placebo analgesia (Wager 2004), contextual-cascade model (Benedetti 2014, *Neuron*); open-label placebo works even when disclosed (Kaptchuk 2010); neurocinematics (Hasson 2008, *Projections*); cinematherapy scoping review (Sacilotto 2022, *Front. Psychol.*). Full linked bibliography in the verification report.

The story happens on your own screen — because that's where your life happens.

Screenlife is the format Timur Bekmambetov pioneered in cinema (*Unfriended*, *Searching*, *Profile*). Content Therapy makes it therapeutic: the narrative arrives as messages, calls, notifications and video woven into the patient's actual device — not a chatbot bubble, not an app you forget to open.

Why it beats “detached AI”

The #1 documented user complaint about AI mental-health apps is detachment from real life. Screenlife's given-circumstances construction builds the story out of the user's actual context — their schedule, their constraints, their world.

Engine specifications

Latency-aware dispatch (<2s channel decisions) · coherence validation (<1% cross-modal contradictions) · emotional state persists across sessions — quit mid-scene, resume 6 hours later, the story remembers.

Telemetry, not passive viewing

Response latency, choice patterns, task acceptance — rich interaction telemetry closes the adaptation loop in <24h. This is the moat over static branching apps (Betwixt) and conversation-reset chatbots.

Format defensibility: nobody else combines adaptive narrative generation + clinical telemetry + a device-native screenlife canvas. Existing narrative apps are pre-written branching; existing DTx are worksheets with progress bars.

Celebrate the shower. Never punish the streak.

Our user research surfaced one finding that shapes the whole product: depressive and anxious users are *harmed* by conventional gamification. Streaks, overdue badges and task counts create a shame spiral — “on my worst days, showering is an accomplishment.”

Design consequence

- ▶ Micro-accomplishments become plot-critical narrative beats — characters validate effort regardless of outcome
- ▶ Real-world actions are framed as story progression, not homework
- ▶ No streak loss, no guilt mechanics — the anti-Todoist for the productivity-app refugees

The paradox we must beat

AI therapy is criticized for enabling avoidance — engagement without behavior change. An entertaining narrative could make avoidance *more* pleasant. So our gate metric is behavioral, not attentional:

≥60%

of pilot users (n=20) must complete a suggested real-world micro-task within 48h of the narrative prompt — this gates all clinical investment

Industry boundary: Lumosity paid a \$2M FTC settlement for outcome claims without RCT-grade evidence. Our rule — engagement language is free, outcome language requires trial data. [\[FTC\]](#)

The Method, computed.

Stanislavsky's system is the oldest rigorous framework for engineering emotional truth. The Conductor Engine implements it as an orchestration layer — six elements, adapted for therapy:

STANISLAVSKY ELEMENT	THERAPEUTIC IMPLEMENTATION
Super-objective	The therapeutic goal from the doctor's brief — e.g. relieve muscle tension, activate parasympathetic response
Given circumstances	The patient's personalized world: safe place, familiar images, sounds, colors — constraint-satisfaction construction
Through-action	Narrative arc from tension to resolution (pain → peace), sequenced on a directed beat-graph
Emotional memory	Personal recollections and anchors of calm woven into scenes
Physical action	Guided breathing and real-world micro-tasks through a character-guide
Tempo-rhythm	Audiovisual pacing calibrated to the target physiological state

Template-first, generative-later

V1 clinical content is human-authored modular scenes (120-scene library), parametrically selected and adapted by the engine — not LLM-generated. LLM generation is confined to the consumer surface and post-clearance stages.

Why we're strict about this

FDA treats generative content as higher-risk than templates; NEDA's Tessa chatbot was destroyed by an unsanctioned generative layer; Character.AI's output was ruled a *product*, not speech (May 2025). Change control over content is a survival requirement, not bureaucracy.

Four phases, each gated by a kill-criterion.

PHASE 1 · M0 0–3 Prove activation 3-scene screenlife prototype + behavioral telemetry (n=20). Gate: ≥60% 48-h task completion. Onboarding <90s vs ChatGPT benchmark. 10-min ARG teaser for studio pitches.	PHASE 2 · M0 3–9 Revenue + insomnia arc First paid ARG delivery (\$150K+ contract). 7-night insomnia narrative (CBT-I stimulus control as detective mystery), n=20 pilot. Parameterized template system.	PHASE 3 · M0 9–15 Clinical feasibility n=40 open-label, 6-week protocol, adjunctive, recruited via sleep clinics. Advisory board formed <i>before</i> content generation. FDA pre-submission prep.	PHASE 4 · SERIES A RCT + payers Powered RCT n=300–400 (Rejoyn precedent: 386 participants, 13 weeks). Parallel payer/employer pilots — clearance is armor, never the revenue plan.
--	---	---	---

Sequencing logic: the riskiest assumption (narrative drives action, not avoidance) is tested first and cheapest. Entertainment revenue de-risks the clinical timeline instead of competing with it.

Three audiences, one stuck human at the center.

The user

Sarah, 29 — employed, medication-averse, mild-to-moderate low mood. Tried journaling (too effortful), ChatGPT (detached), BetterHelp (waitlist + cost), Woebot (quit in a week). Needs something engaging enough to pull her in and concrete enough to push her to act.

The insomnia crisis segment — third consecutive sleepless night, terrified of psychotropic meds, urgently searching for a structured non-drug program.

The buyers

- ▶ Studios & brands — ARG campaigns for franchise launches (the market that already runs Severance-style activations)
- ▶ Consumers — \$15–20/mo emotional-fitness subscription
- ▶ Employers/EAPs — wellness pilots at \$50–100/user/yr, the proven bridge while payer coverage matures
- ▶ Clinicians & payers — post-clearance, via CMS G-codes

The job to be done

“Get me to actually do the small thing I know I should do — text a friend, go outside, sleep — instead of just thinking about why I can’t.”

Not therapy replacement. We fill the six days between sessions, and serve the majority who take no action at all.

\$2.5M seed, built bottom-up. Revenue starts month 6, not month 36.

Seed allocation — with the reasoning

- ▶ 45% · \$1.125M — engine + consumer product: ~8 FTE lean/distributed team × 18 months. The engine is parameterized orchestration over templates, not frontier-model training — a seed-sized build.
- ▶ 20% · \$500K — content: 120-scene modular screenlife library + ARG teaser. Screenlife is the cheapest professional format that exists: *Unfriended*, a full theatrical feature, was produced for a reported ~\$1M — Bazelevs knows this cost curve better than anyone.
- ▶ 20% · \$500K — clinical: advisory board + protocol + n=40 feasibility. Benchmark: eQuoo ran a published RCT and an NHS listing on £1.5M of *total* company funding. [UKTN]
- ▶ 15% · \$375K — go-to-market: consumer launch (organic-first, Finch playbook) + employer pilot sales.

Revenue build across the seed

- ▶ Months 0–3: pre-sold pilot ARG at \$50–75K — willingness-to-pay test that pays for itself
- ▶ Months 4–9: first full ARG contract \$150K+; target 2 campaigns in year 1
- ▶ Months 6–9: employer/EAP letters of intent at \$50–100/user/yr — contractual revenue independent of insurance
- ▶ Months 9–12: consumer subscription launch (\$19.99/mo, 7-day trial)
- ▶ Post-Series A: PDT on the CMS G0552–G0554 rails — the payment codes Pear died waiting for now exist (effective Jan 2025)

Worst-case runway math: \$2.5M with zero revenue ≈ 18 months at ~\$140K/mo burn; each ARG contract adds ~1+ month of runway.

Deliberately flagged: \$2.5M began as an internal estimate (our own pack manifest marks it unsourced) — the bottom-up build above is our justification, and the two riskiest inputs (ARG fees, consumer WTP) have pre-committed tests before the money depends on them.

The model, with its evidence status on every line.

INPUT	VALUE	STATUS	JUSTIFICATION / HOW WE TEST IT
Consumer price	\$19.99/mo	ASSUMPTION	Anchors: therapy ~\$143/session (verified, Headway), ChatGPT \$20/mo, Headspace ~\$12.99/mo, Calm \$69.99/yr. Narrative-app WTP has no public data — pricing A/B at launch.
Net ARPU after store fees	\$14–17/mo	VERIFIED	Apple/Google published 15–30% commission tiers.
Median paying lifetime	4–6 months	BENCHMARK	Wellness-subscription churn ceiling (Calm/Headspace retrenchments); dCBT-I pooled completion 55.9% [npj 2025]. Cohort-measured months 9–12.
LTV (net)	\$56–102	DERIVED	Net ARPU × lifetime; our narrative-retention bet (eQuoo's low attrition, Finch's stickiness) is what pushes toward the top of the range.
Blended CAC	≤\$35 target	ASSUMPTION	Organic-first (Finch scaled on near-zero paid CAC via TikTok); paid-UA sensitivity \$50–80 breaks the model — kill-criterion: payback >12 months.
LTV / CAC	1.6–2.9×	DERIVED	Viable only with organic engine + B2B mix — which is why D2C is never the only channel.
Gross margin	75–80%	BENCHMARK	Software COGS (hosting, inference, support). Marginal content cost/user → ~\$0 by design: the parameterized template engine <i>is</i> the unit-economics thesis (vs Zombies, Run!'s content treadmill).
ARG campaign fee	\$150–300K	ASSUMPTION	Sized against studio experiential budgets (\$500K–2M for genre campaigns); no public ARG price list exists — tested by the \$50–75K pre-sale, months 0–3. Project margin target 40–60% (studio supplies IP).
Employer pricing	\$50–100/user/yr	BENCHMARK	Consistent with EAP/wellness per-employee pricing tier (Calm Business / Headspace for Work class); validated via LOIs months 6–9.
Scale scenario	100K subs → ~\$24M gross / ~\$17–20M net ARR	DERIVED	~70%+ contribution margin at scale; content library is a fixed cost, not per-user. B2B and PDT revenue stack on the same engine and library.

Honest reading: the model works at the top of the ranges and breaks at the bottom — which is exactly why Phase 1–2 test the three ASSUMPTION rows (price, CAC, ARG fee) before they carry any weight.

18 months of seed to Series A.

Evidence compounding for years after.

HORIZON	MILESTONE	PROOF PRODUCED
Month 3	Activation gate passed / failed	≥60% 48-h task completion, onboarding <90s
Month 6	First paid ARG shipped	\$150K+ contract, brand case study, engine works in production
Month 12	Consumer app live + insomnia pilot done	Retention cohorts, ≥15-min sleep-onset improvement in 50%+ of n=20
Month 15–18	n=40 feasibility + FDA pre-sub meeting	Series A package: revenue + users + clinical signal
Series A +13 wks trial	Powered RCT (n=300–400)	Rejoyn precedent: 386 participants, 13 weeks, cleared Apr 2024

The 8-week supervised EMDR-informed pilot (trauma track) runs as a separate expert-review stream and never blocks the commercial timeline.

Screenlife's inventors + a clinically honest bench.

Timur Bekmambetov · Bazelevs

Creator of the screenlife format (*Unfriended*, *Searching*, *Profile*) — the only team that has shipped device-native narrative at box-office scale. Studio relationships feed the ARG pipeline directly.

Stanislavsky AI

The engine team: adaptive narrative generation, psychological profiling, platform-scale content delivery. Patent refiling underway for latency-aware dispatch and emotional-state coherence validation.

Clinical bench

Independent clinical revision already commissioned (V. Kravitz, May 2026) — it reshaped our positioning before a single investor saw it. Advisory board (trauma specialists, sleep-medicine MDs, CBT-I) forms before content generation, with veto power over consumer features.

Honest gaps we're hiring for: a clinical lead with DTx trial experience, and independent (non-founder-authored) study PIs — the category's evidence is almost entirely company-affiliated, and independent replication is the unclaimed differentiator.

It's March 2028. Content Therapy is dead. What killed it?

Before this pitch we ran a formal pre-mortem and mapped every failure story against the documented graveyard of this category. Top findings:

FAILURE STORY	LIKELIHOOD	IMPACT	MITIGATION + CHEAPEST TEST
We bet on reimbursement — clearance ≠ coverage (Pear: \$1.6B → \$6M; Akili: \$1B → \$34M)	HIGH	Fatal	Entertainment + employer revenue first; clearance is optional armor. Pre-sell pilot ARG at \$50–75K in months 0–3.
Narrative entertains avoidance — great retention, failed clinical endpoints	MED	Fatal	Phase-1 behavioral gate (≥60% 48-h task completion) before any clinical spend.
Onboarding loses to ChatGPT — >90s profiling → abort to the \$0 default	HIGH	Severe	A/B “play immediately vs full intake,” n=200; hard 90-second ceiling.
Claims overreach — “EMDR encoded in cinema” draws FTC/clinical backlash (Lumosity \$2M)	MED	Severe	Kravitz language rules adopted; every public claim mapped to an evidence tier.
AI safety incident — generative content harms a vulnerable user (Tessa; Character.AI liability)	MED	Fatal	Human-authored clinical content; hard-coded crisis escalation; vendor change-control clauses; age assurance.
Consumer economics grind us down — CAC/churn (Calm –20% staff, Headspace ×3 layoffs)	HIGH	Severe	B2B channels from month 6; D2C treated as brand-building, modeled with honest churn.
Privacy monetization reflex — one ad pixel on intake = FTC (BetterHelp \$7.8M, Cerebral \$7M)	LOW	Severe	Zero third-party trackers on any authenticated surface, from day one, audited.
Content treadmill — narrative production outruns revenue (Zombies, Run! near-death)	MED	Moderate	Parameterized template reuse; ARG content is studio-supplied IP.

Full 12-risk table with early-warning signals and sources — in the risk report accompanying this deck.

Every number in this deck survived a fact-check. These didn't.

We audited our own draft materials against primary sources and corrected them — the corrections are part of the pitch.

DRAFT CLAIM	VERIFIED REALITY
"\$5.5B market (2024) → \$19B by 2030"	No major firm supports it. Grand View: AI-in-mental-health ~\$1.7B (2025) → \$5.1B (2030); mental-health apps → \$17.5B by 2030. We use the real numbers.
"\$265 average therapy session"	~\$143 average (Headway, 175K+ therapists); \$100–200 typical; \$265 is a big-city specialist ceiling.
"42% of therapists take no insurance"	~1 in 3 psychologists (APA Practitioner Pulse 2024); ~half of psychiatrists.
"Rejoyn trial: 389 participants"	386 participants, 13-week Mirai trial; cleared Apr 2, 2024.
"CMS proposed codes (2024)"	Stronger than draft: codes finalized — G0552–G0554 effective Jan 1, 2025; expanded 2026.
"Digital CBT-I completion 45–67%"	55.9% pooled (npj Digital Medicine 2025 meta-analysis, range 17–86%).
"Wysa FDA Breakthrough, Apr 2026" · "Youper \$14.99/mo"	Wysa BDD = May 2022; Youper is \$69.99/yr and shuts down Sep 30, 2026 — removed as comparable.
"432 Hz + binaural beats" (protocol spec)	Overstated. Modest meta-analytic support for auditory beats on anxiety; 432 Hz = small pilots only. Framed as experimental design choice, not evidence.

Key sources: [Kam-Hansen 2014](#) (migraine placebo) · [Hasson 2008](#) (neurocinematics) · [Kaptchuk 2010](#) (open-label placebo) · [CMS CY2026 PFS](#) · [Otsuka/Rejoyn clearance](#) · [Gallup depression 2026](#) · [npj Digital Medicine 2025](#) (dCBT-I) · full linked verification report + market intelligence in the shared Drive folder.

Everything behind this deck, one click away.

The full pitch pack lives in the shared Google Drive folder; the two companion presentations are live web artifacts.

Google Drive — Pitch Pack folder

- ▶ 📁 Pitch Pack – Doronichev & Garber (Sept 2026) — the folder with everything below
- ▶ 00 · README – Pitch Pack Index
- ▶ 01 · Pitch Deck – Full Content (EN)
- ▶ 02 · Data Verification Report — 24 claims checked, 8 corrected, all sourced
- ▶ 03 · Pre-mortem Risk Register — 13 risks + kill-criteria

Research & companions

- ▶ 04 · Market Intelligence – 14 Players
- ▶ 05 · Failures & Lawsuits Casebook — 12 cases, amounts & dates
- ▶ 07 · Business Plan & Unit Economics – assumptions justified
- ▶ 🗺️ Competitor Atlas — 20 slides, one per competitor (incl. the Лоргопотам verdict)
- ▶ 📊 Market Intelligence dossier — full tables + risk register, interactive

Access note: the Drive documents open for accounts the folder is shared with; the two artifacts are private until shared from their page menus. Source materials (Content Therapy Bible, EMDR clinical revision, full_pack export) remain in the parent Drive folder.